

The background is an abstract composition of several large, overlapping triangles in various colors: red, orange, yellow, teal, blue, and purple. The triangles are separated by thin white lines, creating a dynamic, geometric pattern.

LINUXCONFAU
ONLINE 2021

Large Virtual Address support (52-bit) in ARM64 kernel

SPEAKER

Bhupesh Sharma

Red Hat

<bhupesh.linux@gmail.com>

Twitter: [@bhupesh_sharma](https://twitter.com/bhupesh_sharma)

\$whoami



- Part of Red Hat kernel team.
- Been hacking on boot loaders & kernel since past 15 years.
- Contribute to:
 - Linux,
 - EFI/u-boot bootloader, and
 - User-space utilities like:
 - kexec-tools, and
 - makedumpfile.
- Co-maintain crash-utility tool




Outline

- Large VA support for arm64
 - What?
 - How?
- 52-bit VA kernel support - arm64
- Flipping the arm64 kernel memory layout
- Impact on user-space applications
- 52-bit userspace VAs
- Next Steps

Large VA support for arm64 – What?

- 64-bit hardware ➡ can address **very large** address space.
 - Upto 16 **EiBs** ($16 \times 1024^6 = 2^{64} = 18,446,744,073,709,551,616$ bytes)
 - Approx 18.4 **exabytes** of memory. 
- Servers available with ➡ **64 TiB** (& upwards) of memory.
 -  Use-cases ➡ requiring ➡ addressing spaces $> 2^{48}$ bytes.
- **Limitations**
 - Not all instruction sets, and
 - Not all processors,
 - support a full 64-bit virtual or physical address space.




Large VA support for arm64 – What?

- x86_64

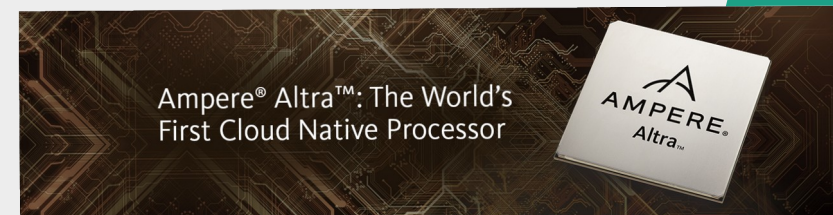
Reference:

- Supports 5-level page tables in both Hardware & Software.
- Allows addressing address space = 2^{57} bytes.
- Bumps limits to
 - → 128 PiB of virtual address space,
 - → 4 PiB of physical address space.



- arm64

- Introduces 2 new architecture extensions
 - 52-bit addressing extensions
 - ARMv8.2 LVA, and
 - ARMv8.2 LPA
- Allows addressing
 - → 4 PiB of virtual address space,
 - → 4 PiB of physical address space.



Large VA support for arm64 – How?

- **ARMv8.2 LVA**

- Supports larger VA space
- Each translation table base register of up to 52 bits
 - when using the 64KB translation granule.

- **ARMv8.2 LPA**

- Allows larger intermediate physical address (IPA), and
- PA space of up to 52 bits when using the 64KB translation granule.
- Allows a level 1 block size where the block covers a 4TB address range for the 64KB translation granule if the implementation support 52 bits of PA.

NOTE: These features are supported in [AArch64](#) state only.

- Cortex-A processors with ARMv8.2 extension support:

- Cortex A55
- Cortex A75
- Cortex A76

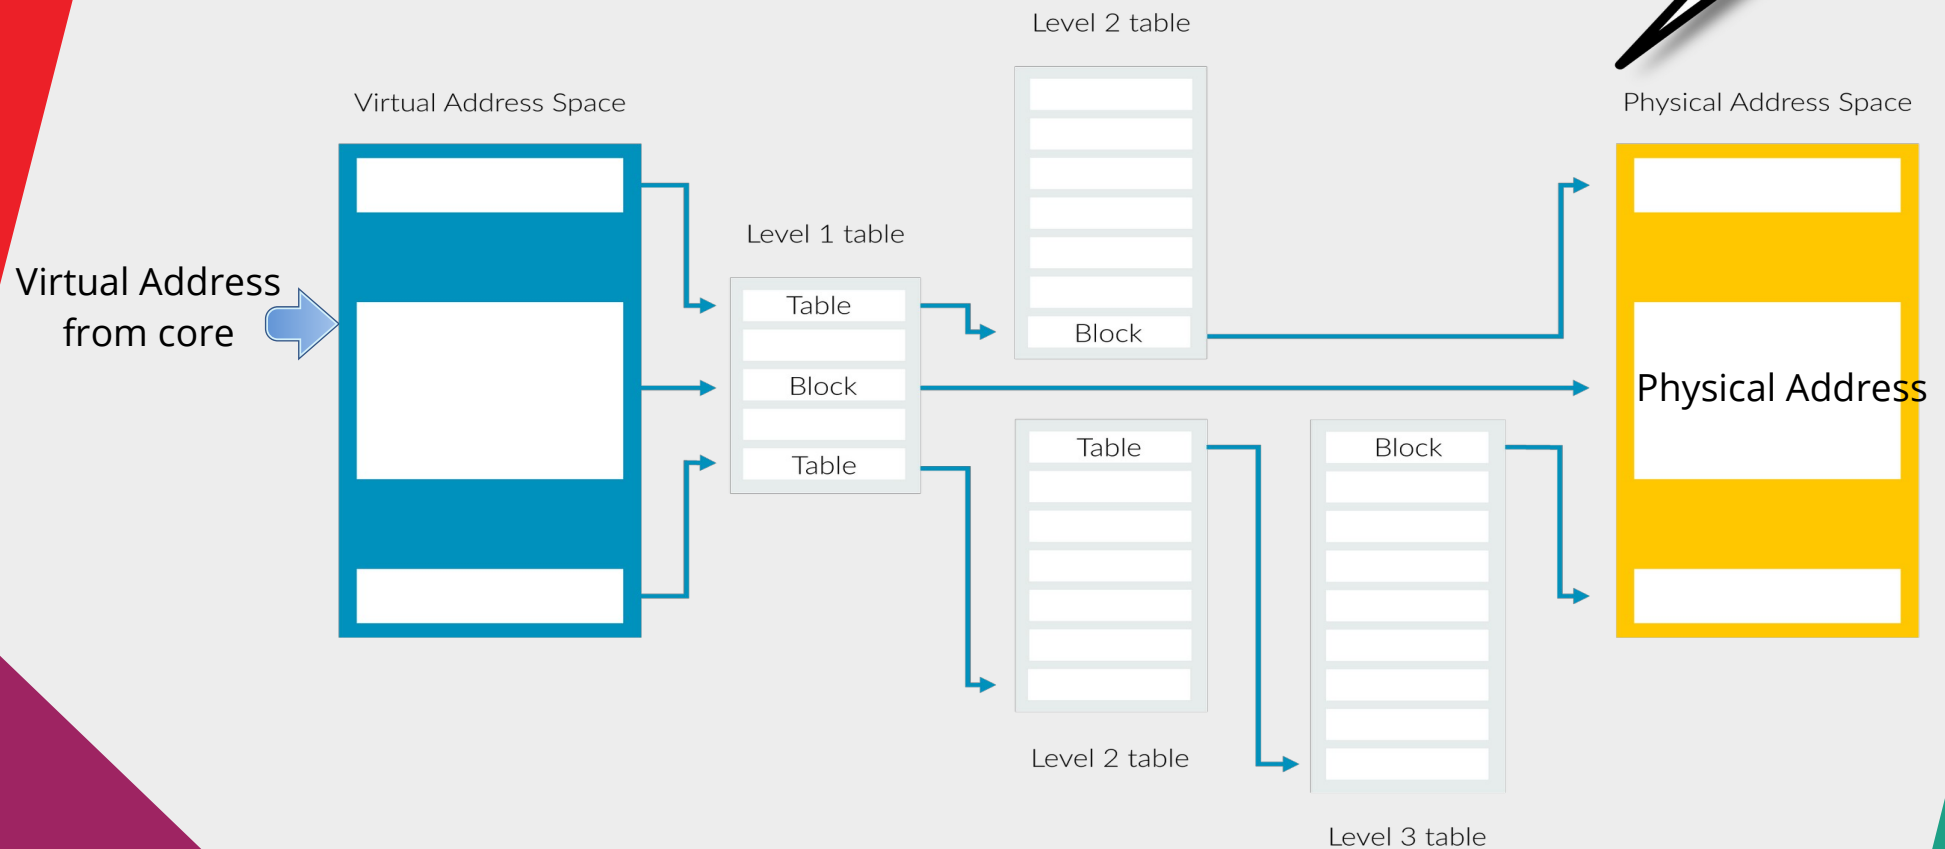


Hardware
Support

Large VA support for arm64 – How?

- Translating a virtual address to a physical address

Hardware
Support



Reference: [Memory management guide](#) from ARM

Large VA support for arm64 – How?


- arm64 Linux uses:

- 4KB page size

- 39-bit (512GB) virtual addresses \Rightarrow 3 translation tables levels
 - 48-bit (256TB) virtual addresses \Rightarrow 4 translation tables levels.

Software
Support

Translation table lookup with 4KB pages:



Reference:

[Documentation/arm64/memory.rst](#)

Large VA support for arm64 – How?

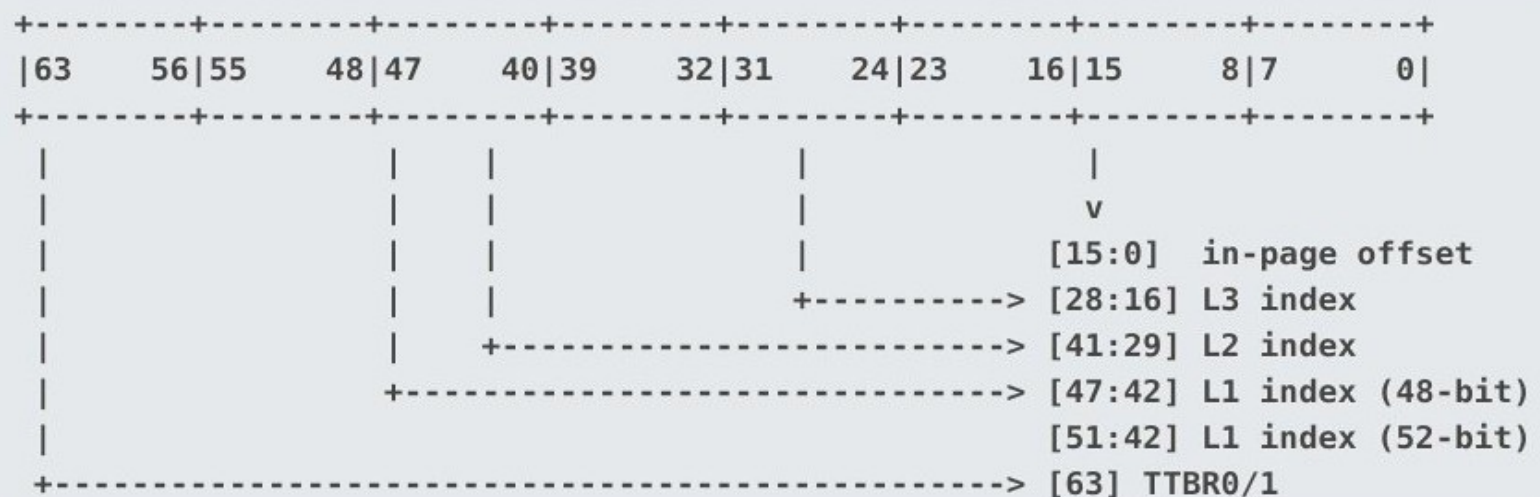
- arm64 Linux uses:
 - 64KB page size
 - 42-bit (4TB) virtual addresses \Rightarrow 2 translation tables levels,
 - but the memory layout is the same.
 - 48-bit (256TB) virtual addresses \Rightarrow 3 translation tables levels
 - 52-bit (4PiB) virtual addresses \Rightarrow 3 translation tables levels with ARMv8.2 extension
 - expands number of descriptors in the first level of translation.

Software
Support

Reference:


[Documentation/arm64/memory.rst](#)

Translation table lookup with 64KB pages:



Large VA support for arm64 – How?

- A **sample** arm64 translation table walk




52-bit VA kernel support - arm64

- **Design problem** from a software support p-o-v
 - Older arm64 CPUs which **don't support** ARMv8.2 extensions.
 - New / Upcoming arm64 CPUs **which support** ARMv8.2 extensions.
- **Selected design** approach
 - Have a **single** kernel binary
 - At **early boot time** check if the ARMv8.2 hardware feature is present or not.

Reference:

[Documentation/arm64/memory.rst](#)



52-bit VA kernel support - arm64

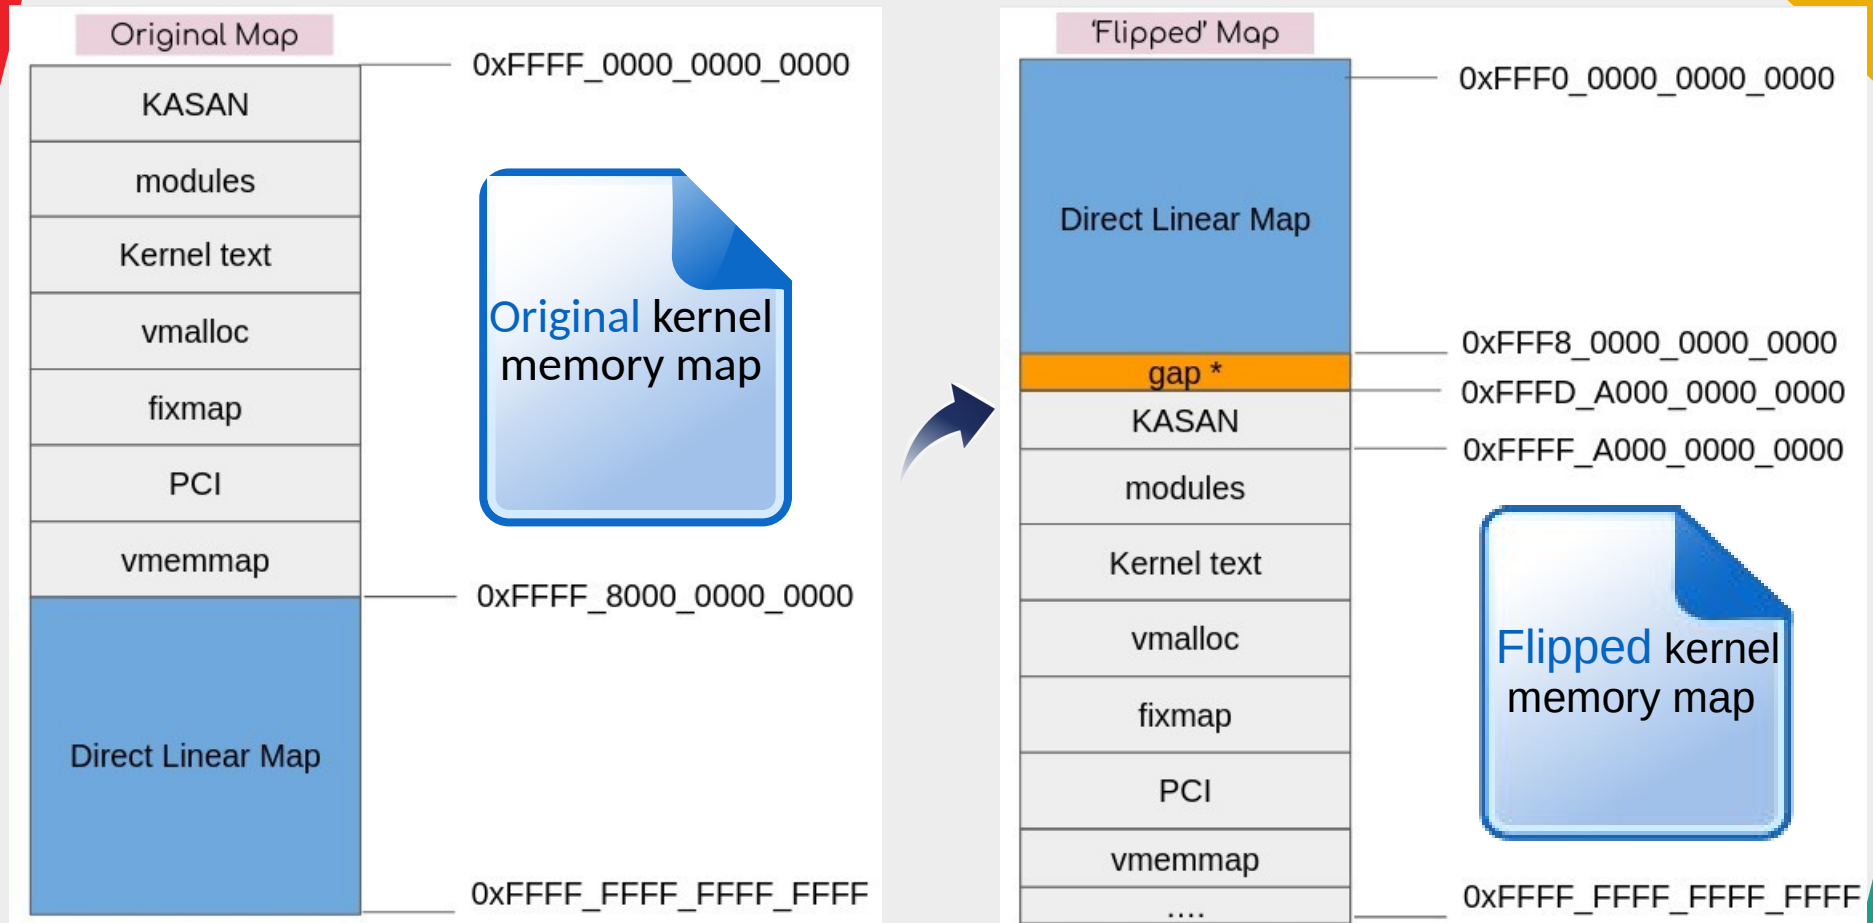
- **Single kernel binary** for both 48-bit and 52-bit VA spaces.
- **VMEMMAP** constraints
 - must be sized large enough for **52-bit VAs**, and
 - must be sized large enough to accommodate a fixed **PAGE_OFFSET**.
- VA bits related **variables** used by kernel code:

VA_BITS	Compile time constant	Maximum size of VA space, used for things like static array and region sizes
VA_BITS_MIN	Compile time constant	Minimum size of VA space, used to ensure pointers are addressable
VA_BITS_ACTUAL	Variable	The actual size of the kernel VA space

* `vabits_actual`

Reference: <Documentation/arm64/memory.rst>

Flipping the kernel memory layout - arm64



Direct Linear Map is in **Higher Half** of the VA space.


Reference: [Documentation/arm64/memory.rst](https://lwn.net/Articles/711111)


- kernel **text** addresses is kept constant, even for 48 to 52-bits migration.
- We need to **flip** the VA space.

Impact on user-space applications - arm64

- User-space applications **impacted** due to **flipped** kernel memory layout
 - used to debug running / live kernels,
 - to analyze vmcore dumps.
 - for example: [kexec-tools](#), [makedumpfile](#) & [crash-utility](#).
- Debugging applications need to perform a [va_to_pa\(\)](#) conversion
 - Walk the translation table(s) for determining the physical address
- These applications are broken upstream currently.
- Have proposed fixes for affected applications
 - some have been **accepted** upstream,
 - others are still **pending**
 - makedumpfile [fix](#), kexec-tools [fix](#)

52-bit userspace VA - arm64



-  What happens to *other existing* applications?
 - To maintain *backward compatibility*
 - *kernel* will, *by default*, return virtual addresses to userspace from a *48-bit* range.
 - *Opt-in* model for willing user-space application(s)
 - Hint parameter is passed to *mmap()* calls to receive addresses in 52-bit range.
- ```
maybe_high_address = mmap(~0UL, size, prot, flags,...);
```
- How to build kernel which returns addresses in 52-bit range?
    - Enable CONFIG options:
      - *CONFIG\_EXPERT && CONFIG\_ARM64\_FORCE\_52BIT*

*NOTE:* Only intended for debugging + should **not** be used in production.



# Next Steps

- Fix broken **userspace** applications – *WIP*.
- Create awareness about the **flipped** kernel address map.
- JIT and other applications need to be made aware about the **mmap()** hint parameter and how they can receive addresses in 52-bit range.
- While I am recording this talk, Ard Biesheuvel has posted a patchset to **extend** the Linear range for 52-bit configurations:
  - <https://lore.kernel.org/linux-arm-kernel/20201008153602.9467-3-ardb@kernel.org/>
- Test upstream kernel on both old CPUs (48-bit VA) and new CPUs with 52-bit VA)
  - In absence of a real 52-bit HW, you can use [ARMv8 fast model](#) simulator for some quick checks.



Slides can be found on github

*Email: <bhupesh.linux@gmail.com>*

*Twitter: @bhupesh\_sharma*

The talk is licensed under a Creative Commons  
Attribution-ShareAlike 4.0 International License.

The image features a vibrant, abstract geometric design. It is composed of several large, overlapping triangles in various colors: a bright red triangle in the top left, a large orange triangle in the top right, a yellow triangle on the right side, a teal triangle in the bottom right, a dark blue triangle in the bottom left, and a purple triangle on the left side. These triangles are separated by thin white lines. In the center, where the triangles meet, there is a white, irregularly shaped space. Within this white space, the word "Questions?" is written in a simple, dark gray, sans-serif font.

Questions?